Skip to main content :/Rag Search in r/Rag Create

r/Rag • 19d ago
MiserableBug140

4 steps to turn any document corpus into an agent ready knowledge base

Discussion

Most teams building on documents make same mistake. Treat corpus as search problem.

Chunk papers, embed chunks, vector store, call it knowledge base. Works in demos, breaks in production. Returns adjacent context instead of right answer, hallucinates numbers from tables never properly parsed, fails on questions needing reasoning across papers.

Problem isn't retrieval or embeddings or chunk size. Embedded text chunks aren't knowledge base, they're index. Index only as useful as structure underneath.

Reasoning-ready knowledge base is corpus that's been extracted, structured, enriched, organized so agent can navigate like domain expert. Not guessing which chunks semantically similar but understanding what corpus contains, where info lives, how pieces relate.

Transformation involves four things most pipelines skip. Structure preservation so relationships stay intact. Semantic tagging labeling content by meaning not location. Entity resolution unifying different names for same concepts. Relational linking connecting related pieces across documents.

Most RAG pipelines do none of these. Embed chunks, hope similarity search covers gaps. For simple lookup on clean prose mostly works. For research corpora where hard questions require reasoning across structure doesn't work.

Building one needs structure-preserving extraction keeping IMRaD hierarchy, enrichment tagging sections by semantic role and extracting entities, indexing supporting metadata filtering and hierarchical retrieval, agent layer doing precise retrieval and cross-paper reasoning.

Tested agent across 180 NLP papers. Correctly answered 93 percent complex cross-paper queries. The 7 percent needing review surfaced with low-confidence flags not returned as confident wrong answers.

Teams building reliable research agents aren't ones with best embeddings or tuned rerankers. They're ones who invested in transformation layer before calling anything knowledge base.

Anyway figured this useful since most people skip these steps then wonder why their agents hallucinate.

6 3 Share

Join the conversation

Sort by: Best Search Comments

Create

document-corpus-into-a-reasoning-ready-knowledge-base-in-2026/

3 Reply ···

johnmacleod99 • 18d ago

Great post, made my day!. You point to pervasive and hard to solve problems.

1 Reply ···

Reddit Rules Privacy Policy User Agreement Your Privacy Choices Accessibility Reddit, Inc. © 2026. All rights reserved.